

Exercises 02 & 03

Guidelines for implementing components of Lecture 03-06 in Group Project

Module 5: Non-Functional Requirements: Performance Optimization

Objective: Implement monitoring and profiling.

Steps:

1. Performance Monitoring:

- Set up **Prometheus** for monitoring.
 - i. Set up metrics collection (e.g., response times, message throughput, database query latency, CPU, memory, request latencies) in each microservice (gRPC, REST FastAPI/Django, Kafka).
 - ii. Expose metrics endpoints (e.g., `/metrics`) in FastAPI and gRPC services.

Note: “Before you can monitor your services, you need to add instrumentation to their code via one of the Prometheus client libraries. These implement the Prometheus metric types.”

<https://prometheus.io/docs/instrumenting/clientlibs/>

- Use **Grafana** to visualize the metrics.
<https://grafana.com/docs/grafana/latest/getting-started/get-started-grafana-prometheus/>

2. Performance Profiling:

- Use Py-Spy or cProfile to analyze Python microservices.
- Identify bottlenecks and optimize code.

3. Load Testing:

- Use **Locust** (ideal for Python) to simulate load and distributed testing.
 - i. Test the performance of synchronous (gRPC/REST) and asynchronous (Kafka) components under high loads

Note: Another popular load-testing tool in the Finnish ICT industry is k6 if your choice of language is JavaScript. You can use either Locust or K6.

Testing:

- Verify the system's throughput under varying loads.
- Check latency and error rates for each microservice.
- Confirm that metrics are correctly captured and visualized.

Module 6: Non-Functional Requirements: Elasticity and Scalability

Objective: Enable the system to scale dynamically based on load.

Steps:

1. Deploy on Kubernetes

- Containerize each microservice using Docker.
 - i. Step-by-step instructions from Linode on building and deploying microservices using Docker and Docker Compose, covering essential aspects such as creating Dockerfiles, setting up Docker Compose configurations, and testing the microservices:
<https://www.linode.com/docs/guides/deploying-microservices-with-docker>
- Deploy on **Kubernetes** for orchestration.
 - i. Stepwise guide (you need to have containerized microservices first):
<https://www.digitalocean.com/community/tech-talks/how-to-deploy-your-application-or-microservice-as-a-kubernetes-deployment>

2. Horizontal Scaling:

- Set up **Horizontal Pod Autoscaler (HPA)** for scaling based on CPU, memory, or custom metrics (e.g., Kafka lag).
 - i. Official walkthrough on configuring HPA based on CPU utilization:
<https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale-walkthrough>
 - ii. how to implement HPA using custom metrics, specifically focusing on Kafka metrics like consumer lag
<https://sunnykr Gupta.github.io/kubernetes-hpa-autoscaling-with-kafka-metrics.html>
 - iii. how to scale Kafka consumer groups using custom metrics in Kubernetes
<https://developers.redhat.com/articles/2023/12/29/autoscaling-kafka-workloads-custom-metrics-autoscaler>

3. Load Balancing:

- Integrate **NGINX** as a reverse proxy and load balancer.
- Configure NGINX to distribute requests to multiple replicas.

4. Elastic Database Scaling (suggestion from CSC colleague.):

- Use CockroachDB or Cassandra, which support distributed scaling
- Implement database connection pooling for performance.

If you have already chosen another database (like not Cassandra or CockroachDB), please feel free to skip elastic database scaling.

5. Elastic Messaging:

- Configure **Kafka partitions** and brokers to scale the messaging system.
- Use **KRaft mode** for Kafka cluster management

Testing:

- Simulate varying traffic levels using Locust.
 - Validate that Kubernetes scales pods dynamically.
 - Test database performance with concurrent read/write operations.
-

Module 7: Replication and Consistency

Objective: Ensure data availability, fault tolerance, and consistency across distributed components.

Steps:

1. Database Replication:

- Configure MongoDB/Apache Cassandra with replication enabled (e.g., Replica Set for MongoDB or multi-datacenter replication for Cassandra, CockroachDB multi-region deployment).

2. Consistency Management:

- Design your services to handle **CAP theorem trade-offs** based on use case (e.g., eventual consistency for Kafka consumers, strong consistency for critical operations).
 - i. Implement eventual consistency in Cassandra or MongoDB.
 - ii. Demonstrate strong consistency using CockroachDB or MariaDB.

3. Distributed Coordination with etcd:

- Integrate etcd for leader election and service discovery.
- Use etcd's watch mechanism for consistent state updates across microservices.

4. Replication in Messaging

- Enable **Kafka topic replication** across brokers for fault-tolerant message delivery.

5. Testing Fault Tolerance:

- Simulate node failures using Chaos Engineering tool: **Chaos Mesh**.

- i. **Chaos Mesh** is easy to integrate into Kubernetes infrastructure without any changes to deployment logic.
- Verify data integrity and recovery mechanisms.

Another popular tool for chaos engineering in the Finnish ICT sector, Gremlin, is not open source!

6. Conflict Resolution:

- Implement a last-write-wins (LWW) policy or vector clocks to resolve replicated data conflicts.

Testing:

- Simulate node failures to verify data replication and recovery.
- Test read/write operations to ensure consistency.
- Test etcd coordination by introducing artificial network partitions and monitoring service recovery.
- Use **Jepsen** to test distributed consistency under network partitions.

Module 8: Resource Management and Scheduling

Objective: Optimize resource allocation and scheduling for efficient system utilization.

Steps:

1. Dynamic Resource Allocation:

- Use **Kubernetes Resource Quotas** and **LimitRanges** to manage resource allocation for each microservice.
 - i. Configure pod-level resource requests and limits for CPU and memory.
- Use **KEDA** (Kubernetes-based Event Driven Autoscaling) to scale Kafka consumers or other services based on event volume.

2. Scheduling Optimization:

- Implement Kubernetes custom schedulers for critical tasks: Configure **Kubernetes affinity/anti-affinity rules** and **priority-based scheduling**.
- Use Kubernetes **taints and tolerations** to segregate workloads (e.g., database on high-memory nodes).

3. Cluster Autoscaling:

- Enable Kubernetes **Cluster Autoscaler** to dynamically add/remove nodes based on demand.

Testing:

- Deploy services with resource limits and test behavior under constrained resources.
 - Simulate node unavailability and verify service redistribution across nodes.
 - Scale workloads and validate autoscaling at the container and cluster levels.
-

Integration Testing

Objective: Test the fully integrated system for functionality and performance.

Steps:

1. Integrated Deployment:

- Deploy all components (gRPC, FastAPI, Kafka, database) on Kubernetes.
- Use Helm charts for consistent configurations.
- Integrate monitoring tools (Prometheus, Grafana) across all components.
 - i. Measure latency, throughput, and resource utilization under integrated workloads.
 - ii. Use Grafana dashboards to analyze system behavior and identify bottlenecks.
- Ensure proper resource allocation and scaling policies.

2. Workflow Validation:

- Simulate end-to-end workflows combining synchronous and asynchronous communication.
- Include database operations to test replication and consistency.

3. Fault Injection:

- Use **Chaos Mesh** to inject failures (e.g., pod crashes, network delays).
- Observe system behavior and verify recovery.

4. Full Load Test:

- Use Locust to generate realistic traffic across all services.
 - i. Conduct load tests using Locust to simulate hybrid workloads (synchronous + asynchronous communication).
- Monitor system performance and resource usage.

Success Criteria:

- Services remain functional under stress and failures.
- Data consistency is maintained across replicas.

- The system scales dynamically with workload.